

kWIP: The k-mer Weighted Inner Product

Kevin Murray

PhD Candidate, Borevitz Lab, CPEB, ANU

26 October 2015



Overview





Large-scale population genomics


- ▶ Moving from 100s to 1,000s or 10,000s of samples *per PhD!*



Large-scale population genomics

- ▶ Moving from 100s to 1,000s or 10,000s of samples *per PhD!*
- ▶ Efficient algorithms to analyse large-scale genomic data
 - ▶ Reference & alignment free: *less bias, de novo*
 - ▶ Platform/protocol agnostic: *future proof*
 - ▶ Computationally efficient: *not the bottleneck*

Fraction of genome




after Peterson *et al.* [1]

Presenting kWIP

- ▶ k -mer based *de novo* estimate of genetic similarity
- ▶ Produces a distance matrix from raw NGS reads

@D594K1-261.C3L3WACXX-3:1101:2729:2264:1:0
 TCGTGAAGGCAGGAAGTACGCAAGCAAGCAAGCAAGCAAGTATGACGC +
 DADH+HFDFH+HH+HC9GCGHCEGGGJHJHJGGG +
 @D594K1-261.C3L3WACXX-3:1101:2729:2264:1:0
 TCGAGAAAGTCGAGCAATCAACGCCCAACCAACTCATGAGTTCTGATTT +
 FTFHF+BFAC+FHGG+GGGHHHHHHHHHHHHHHHHHHHHHHHHHHHHHH +
 @D594K1-261.C3L3WACXX-3:1101:3208:2295:1:0
 TCGAGTGAAGGAAGGAGTCAGACAGCTGATTAAAGCAATCTGTATTAT +
 FTFHHFH+HH +
 @D594K1-261.C3L3WACXX-3:1101:3208:2295:2:0
 TCGAGATTTTAAACATATTAAGTAATGCTCTGCTCGTAATGACACA +
 FFFFFHH +
 @D594K1-261.C3L3WACXX-3:1101:3632:2456:1:0
 TCGAGCATATCATCATCTTTGTCAGTGATGATCATTTTCTGTCGATT +
 FTFHH +
 @D594K1-261.C3L3WACXX-3:1101:3632:2456:2:0
 TCGAGATTAATATCTCTGTTGTTATGACATCTTTCTAGGAAGTAGTA +
 FTFHH






Why Estimate Similarity?

- ▶ Rough approximation of sample relatedness required
 - ▶ For natural collections
 - ▶ As a technical control



Why Estimate Similarity?

- ▶ Rough approximation of sample relatedness required
 - ▶ For natural collections
 - ▶ As a technical control




after Brachi *et al.* [2]



Why Estimate Similarity?

- ▶ Rough approximation of sample relatedness required
 - ▶ For natural collections
 - ▶ As a technical control




after Brachi *et al.* [2]



Why Estimate Similarity?

- ▶ Rough approximation of sample relatedness required
 - ▶ For natural collections
 - ▶ As a technical control




after Brachi *et al.* [2]



Why Estimate Similarity?


- ▶ Rough approximation of sample relatedness required
 - ▶ For natural collections
 - ▶ As a technical control





Why Estimate Similarity?

- ▶ Rough approximation of sample relatedness required
 - ▶ For natural collections
 - ▶ As a technical control





Technological Overview

- ▶ k -mer counting (bag-of-words)
 - ▶ Decompose sequences to overlapping words of k



Technological Overview

- ▶ k -mer counting (bag-of-words)
 - ▶ Decompose sequences to overlapping words of k
- ▶ Hashing and Probabilistic Data Structures
 - ▶ Efficient storage & compute of “bag of words”



Technological Overview

- ▶ k -mer counting (bag-of-words)
 - ▶ Decompose sequences to overlapping words of k
- ▶ Hashing and Probabilistic Data Structures
 - ▶ Efficient storage & compute of “bag of words”
- ▶ (Weighted) Inner Products
 - ▶ Similarity metric weighted by Shannon entropy



kWIP Algorithm

- ▶ For each run: count all k -mers into a Hash
- ▶ For each analysis:
 - ▶ Calculate the entropy of population frequency (H)
 - ▶ For each pair of runs A and B , calculate

$$\sum_{i=0}^n A_i \cdot B_i \cdot H_i$$

A	0	2	0	2	0	1	0	1	0
					•				
B	0	2	1	7	0	1	0	0	0
					•				
H	h_1								h_n



- ▶ The software:
 - ▶ C++, >2000 lines of code
 - ▶ Uses **khmer** for *k*-mer counting & hashing
 - ▶ Parallelised, fast
 - ▶ GNU GPL licensed, source code on GitHub




Demonstration

- ▶ Set of 96 rice runs from 16 samples (6 tech reps ea)
- ▶ Expectations:
 - ▶ All runs cluster into groups of 6 reps (16 samples)
 - ▶ Big split between Indica & Japonica: (7 and 9 respectively here)
- ▶ We see this with kWIP, not with Unweighted IP
- ▶ Took 8 hours on 16 CPU Raijin node, 60-80GB RAM




WIP





IP





Conclusions

- ▶ kWIP is implemented, β software
- ▶ A few users in the audience (thanks!)
- ▶ Further simulations and experiments required
- ▶ Publication in preparation





Thanks

- ▶ Supervisors: J Borevitz, S Forêt, G Huttley and B Pogson
- ▶ Christfried Webers, Cheng Soon Ong, Norman Warthmann
- ▶ **khmer** folks: C. Titus Brown, Michael Crusoe, Camille Scott (DIB-lab) @ UC Davis
- ▶ Beta testers
- ▶ Yourself



References

-  Peterson, B. K. *et al.* Double Digest RADseq: An Inexpensive Method for De Novo SNP Discovery and Genotyping in Model and Non-Model Species. *PLoS ONE* **7**, e37135 (2012).
-  Brachi, B., Morris, G. P. & Borevitz, J. O. Genome-wide association studies in plants: the missing heritability is in the field. *Genome biology* **12**, 232 (2011).

.